Pankayaraj Pathmanathan

Curriculum Vitae

Website | p.pankayaraj.gmail.com | LinkedIn | Github | Full Resume

WORK EXPERIENCE (SELECTED)

CURRENT, FROM SEPT 2022 (FULL TIME)

Research Assistant, Teaching Assistant
University of Maryland College Park

During this time, I primarily worked on LLM robustness, including RLHF poisoning, backdoor poisoning, and reasoning robustness.

MAY 2025 – AUG 2025 AND NOV 2025 - CURRENT (FULL TIME & PART TIME)

Machine Learning Research Intern
Netflix, Los Gatos

Worked on pretraining large language model based embeddings models

FEB 2022 – AUG 2022 (FULL TIME)

Research Engineer
Singapore Management University

Worked on constraint reinforcement learning. This work was **published at AAAI 2023**

PUBLICATIONS (SELECTED)

Pankayaraj P, Huang, F (2025). Teach a Reward Model to Correct Itself: Reward Guided Adversarial Failure Discovery for Robust Reward Modeling [**Oral**] ***ACL 2026*** The 64th Annual Meeting of the Association for Computational Linguistics, California

Pankayaraj P, Sehwag, U. M., Panaitescu-Liess, M.-A., Cho-Yu Jason Chiang, Huang, F. (2025a). Advbdgen: Adversarially fortified prompt-specific fuzzy backdoor generator against llm alignment [**Oral**] **4.6%** , in the **In 40th *AAAI - AIA*** Conference on Artificial Intelligence, Singapore

Pankayaraj P, Chakraborty, S., Liu, X., Liang, Y., Huang, F. (2024). Is poisoning a real threat to LLM alignment? maybe more so than you think, In [Poster], In 39th ***AAAI - AIA*** Conference on Artificial Intelligence Philadelphia, Pennsylvania, USA

Pankayaraj P, Varakantham, P. (2022). Constrained reinforcement learning in hard exploration problems [Poster], In 37th ***AAAI*** Conference on Artificial Intelligence Washington, D.C. USA

Panaitescu-Liess, M.-A., **Pankayaraj P**, Y. K., Che, Z., An, B., Zhu, S., Agrawal, A., Huang, F. (2024). Poisonedparrot: Subtle data poisoning attacks to elicit copyright-infringing content from large language models [**Oral**] , in the ***NAACL*** 2025

Pankayraj P, Rodríguez, N. D., Ser, J. D. (2023). Using curiosity for an even representation of tasks in continual offline reinforcement learning, In ***Cognitive Computation*** Journal 2023

Panaitescu-Liess, M.-A., Che, Z., An, B., Xu, Y., **Pankayaraj P**, Chakraborty, S., Zhu, S., Goldstein, T., Huang, F. (2024). Can watermarking large language models prevent copyrighted text generation and hide training data?, In 39th ***AAAI*** Conference on Artificial Intelligence Philadelphia, Pennsylvania, USA

Pankayaraj. P, Maithripala, D. H. S. (2020) A decentralized communication policy for multi agent multi armed bandit problems [**Oral**], In ***European Control Conference*** 2020, Saint Petersburg, Russia

Pankayaraj P, Maithripala, D. H. S., Berg, J. M. (2020). A decentralized policy with logarithmic regret for a class of multi-agent multi-armed bandit problems with option unavailability constraints and stochastic communication protocols [**Oral**], In 59th ***IEEE Conference on Decision*** and Control, Jeju Island, Republic of Korea

EDUCATION

CURRENT	**PhD computer science** ADVISOR: FURONG HUANG University of Maryland College Park.
2015-2020	**BSc Computer Science** UNIVERSITY OF PERADENIYA, SRI LANKA

REFERENCES

NAME	**Prof. Furong Huang**
EMPLOYER	University of Maryland College Park
NAME	**Ashish Rastogi**
EMPLOYER	Netflix
NAME	**Hafez Asgharzadeh**
EMPLOYER	Netflix
NAME	**Prof. Pradeep Varakantham**
EMPLOYER	Singapore Management University

WORKSHOPS (SELECTED)

Panaitescu-Liess, M.-A., Che, Z., An, B., Xu, Y., **Pankayaraj P**, Chakraborty, S., Zhu, S., Goldstein, T., Huang, F. (2024). Can watermarking large language models prevent copyrighted text generation and hide training data?, In [***Best Paper***] ***NeurIPS Workshop AdvML-Frontiers*** 2024

Pankayaraj P, Sumanasekera, Y., Samarasinghe, C., Elkaduwe, D., Jayasinghe, U.,Maithripala, D. H. S. (2020). Multi-agent reinforcement learning in sparsely connected cooperative environments[***Best Research Paper***], in ES-CaPe 2020, Sri Lanka.

Pankayaraj P, Huang, F (2025). RAGPart & RAGMask: Retrieval-Stage Defenses Against Corpus Poisoning in Retrieval-Augmented Generation ***AAAI Workshop on New Frontiers in Information Retrieval*** 2026 [**Oral**]

Pankayaraj P, Huang, F (2026). RAGPart & Deliberative Alignment is Deep, but Uncertainty Remains: Inference time safety improvement in reasoning via attribution of unsafe behavior to base model ***COLM Actionable Interpretability Workshop*** 2026